Contents

List of Figures

Chapter 1

Introduction

Upon starting their studies in computer science, new students may find the concepts of programming in general, and object oriented programming in particular, to be difficult to grasp. The understanding of computers, and how they work, is not given a lot of focus in the school system, and only a few students have actually tried learning programming before coming to the university. With little or no previous knowledge, it can be hard to understand the concepts, and to get a mental model of what is going on when writing a program. Especially code not written by themselves, for example exercise frameworks, and code generated by various tools, can be hard to get a good understanding of. Traditional debuggers are not necessarily helping when detecting a runtime error, and often significant amounts of time is spent searching for the cause of a bug, instead of actually fixing it [Ko and Myers, 2006]. Tools that present the state of a program in a simple, visual way may help to understand how a program works, and how its components interact with each other.

During the second year of the computer science, and the informatics studies at NTNU, there is an increased focus on projects and more complex software. Among the mandatory courses of these studies, we find the course Man-machine-interaction (MMI). This course handles topics related to creating graphical interfaces, and specifically how to implement a *Model View Controller* (MVC) architecture in the Java Swing framework. The learning goals for the MMI-course are as follows:

> Introduce the student to concepts, methods and techniques for designing man-machine-interfaces, knowledge and skills in object oriented construction of graphical window-based interfaces.

Through the course, students are introduced both to the theories of good user interface design, and how such interfaces can be implemented in Java. Java, and the *Integrated Development Environment* (IDE) Eclipse, make a common framework for many of the programming courses at NTNU, making them the main focus when

looking for tools to support the students.

1.1   Research Questions

The goal of this report is to look into the available tools that can aid developers, focusing on the potential use in the MMI-course. This implies a focus on how to better understand the interactions within programs implementing the MVC pattern, as well as the behavior of graphical interfaces. Looking into the tools, it will identify the different types, as well as their strengths and weaknesses, using this as a basis for comparison. After comparing, and determining which tool is likely to be of most use, the report will look at both the potential for improvements, and how it may fit into the teaching process, including the effect of any improvements made. To be more specific, the following research questions have been formulated:

Research Question 1 *What is the current state of the various visualization tools that are available?*

Research Question 2 *Could any of these be integrated into the current teaching environment at NTNU, consisting of Java and Eclipse?*

Research Question 3 *Is there room for improvement in how these tools are used, and the ease of using them?*

Research Question 4 *Would the use of such tools and any improvements actually be useful for the students?*

Chapter 2

Task Description

In order to answer the research questions, the report is divided into appropriate sections. The prestudy in chapter 3 will attempt to gain an overview of the different methods used by the tools in question, as well as tools that implement them. Some tools that does not fit into the teaching environment, due to language- or IDE-support, are still mentioned because they serve as good examples of the functionality they provide. After gaining an overview, the most fitting tool will be determined through a comparison of functionality and environment will be examined further, detailing its features and identifying areas that may be improved form a usability standpoint.

Not all potential improvements can be expected to be implemented, as it depends on the internal workings of the selected tool. Some cosmetic changes may be easy to implement, while others may require a rewrite of large parts of the system, which ultimately will be considered to expensive, time-wise, compared to the potential benefits.

Determining the actual usefulness of the selected tool will be done through testing with students, mainly those that are in their second year at the time of writing, but depending on the access to volunteers, students at other levels may be used as well.

Chapter 3

Prestudy

This section will serve to get an overview of the current situation of debugging methods, and the various tools that are available.

3.1   Methods

There are several ways a debugger can aid the programmer beyond just showing the current state of a program at a *breakpoint* chosen before running. For a fresh programmer, either in general, or at a certain project, the most useful method is probably to generate diagrams that visualize the current state, and the path of execution. I.e. some form of object-, and/or sequence-diagram. Such diagrams can make it easier to get an overview of a programs current state, see the contents of objects and how they relate to each other, and to understand how the various components work together.

In order to generate the diagrams, the tools can analyze both the source-code, and an *execution trace*, depending on the type of diagrams to generate. For a general overview, a class diagram, showing how object types are composed, can be generated by simply analyzing the source code. But in order to get a more specific view, for example of how the different components interact, an analysis, either in real-time or from an execution trace, of a running program is needed. Such an analysis can be used to generate object diagrams that show how objects interact and what values they have at a given time. One can also get a sequence diagram, showing what methods were called, and in what order.

Execution traces can also be used to enable backwards stepping of program execution. Stepping back in time allows the user to not only see the failure state of a program, but to go back and see what caused the problem, instead of adding a new breakpoint and running the program again. Each reverse step can be fairly cheap, but it may still be impractical to make large jumps in the execution history.

One can avoid the potential disadvantages of manual backstepping by using queries instead. Queries enable the user to asks the debugger about the current and earlier states of execution in a simple way. The debugger then does the work of finding what was asked for, instead of the user manually searching through the program states.

A major downside of execution traces, is the performance penalty from doing extra work for every step in the execution. The amount of work will vary, depending on implementation details and of course what is done with the log. If all the tracing-process does is write to a log-file, which will be used later, one can achieve significantly smaller performance impact compared to a system that does real-time analysis of the data.

3.2 Existing Tools

There currently exists several tools that provide one or more of the methods mentioned above. The following list is not extensive, but includes some of the most interesting tools, considering the focus of this report. Some of the tools listed below does not support Java or eclipse, but provide interesting features that are worth mentioning, despite their obvious inability to be integrated into the desired teaching environment.

GNU debugger (GDB) offers a tracing environment, and support for many languages, but not Java. Due to its *Command Line Interface* (CLI), it is not necessarily easy to use on its own, and so, there are several front-end platforms that provides a graphical environment around GDB, including Eclipse.

Code Canvas uses an interesting way of visualizing an entire project, everything from source-code to design documents and diagrams are layered onto a large canvas, allowing easy navigation between various elements, but is restricted to Microsoft Visual Studio, and the languages it supports.

The *Trace Viewer Plugin* [Kranzlmüller and Klausecker, 2009] for g-eclipse, uses a trace to generate visualizations of the program execution, and thus makes it easier to understand, but is designed for massive parallelism, requires a special version of eclipse, and may not be very useful for understanding smaller programs.

Whyline [Ko and Myers, 2009], and the *Trace-Oriented Debugger* [Pothier et al., 2007] also utilize execution traces, but use them to enable querying, instead of providing visualizations. Additionally, Whyline exists only as a separate application, and does not integrate into any IDE.

The *Debug Visualization Plugin* for eclipse provides an alternative to the variable view provided in eclipses debug-perspective, providing a graph that represents the

variables of a program. The user still needs to use regular techniques they would use normally in order to pause the program-execution, and be able to actually view the state of the variables.

JAVAVIS [Oechsle and Schmitt, 2002] provides visualizations in the form of UML-diagrams, but does not provide any debugging features.

Jinsight [Pauw and Vlissides, 1998] is a powerful tool built by IBM, supporting both tracing and visualization. However, it is restricted to z/OS and linux on system Z, preventing most people from using it.

JIVE is a tool that utilizes execution traces to provide diagrams while running a program. Developed at the university of Buffalo, it is installed a an Eclipse plugin, and provides several new views to display the information it provides. In addition to providing diagrams, the trace log is also used to enable backstepping, which is coupled with the diagrams to always show the selected execution state.

Of the mentioned tools, the Trace Viewer Plugin, Trace-Oriented Debugger, Debug Visualization Plugin, and JIVE are all available as eclipse-plugins, and are thus fairly simple to integrate into the existing teaching process. The features they provide, on the other hand, vary. The Trace Viewer Plugin is as mentioned, designed for massive parallelism, and does require a specific version of eclipse, and because of this, it is not really suited for further study.The Trace-Oriented Debugger provides a debugging environment supported by trace logs, as implied by the name. It presents its information mostly in a textual way, and does not provide any visual diagrams of the program structure, or execution order. The Debug Visualization Plugin expands on the debugging functionality of Eclipse by providing a visual view of variables, and is designed to be used alongside the rest of the debugging environment provided by eclipse.

JIVE seems to be the only tool that utilizes all three methods mentioned in section 3.1, as well as being freely available as a plugin for eclipse, making it easy to install and use. During program execution, Jive generates a *contour diagram* [Jayaraman and Baltus, 1996], and a sequence diagram. Combined with an execution trace, it allows the user to jump back and forth in the execution, and have the diagrams updated accordingly. Querying is supported with pre-defined search-templates added to the built-in search window in Eclipse.

Due to all the extra work being done when using jive to debug a program, the performance is not always acceptable. For small non-interactive programs, the added waiting time may not be a problem, but larger programs may suffer from a significantly longer execution time, and even simple interactive programs can use up to a second to respond to input on a fairly powerful computer.

Chapter 4

Enhancing JIVE

This section will first provide details on the various features of JIVE, and how they are used. This will be followed by an identification of various shortcomings as well as suggestions on how those areas may be improved. Next, the features and improvements that were implemented, and how those implementations are used will be detailed, followed by the results of user-testing with a group of students.

4.1 Features
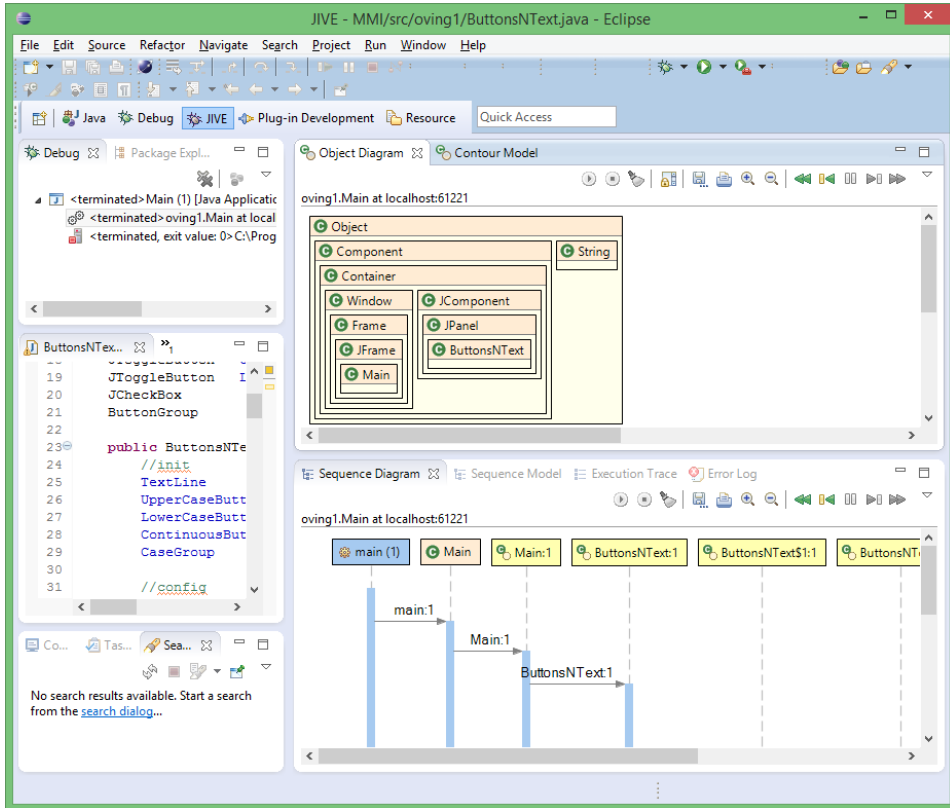
Figure 4.1: The JIVE perspective in Eclipse

As mentioned in the prestudy, JIVE installs as a plugin in the eclipse IDE, adding another perspective in the environment shown in Figure 4.1. The default views shown in this perspective are the "object diagram" and "sequence diagram" views, making the diagrams generated during debugging JIVEs two most apparent features. The diagrams are updated according to the current state of the debugged program, so that the backtracking functionality allows you to see the entire execution graphically step-by-step. The other views provided by JIVE are as follows: "contour model", "sequence model", "execution trace" and "search".

Figure 4.2: A contour diagram generated by JIVE

The object diagram-view shows the current state of the program by using a contour-diagram. Contour-diagrams, as shown in Figure 4.2, are based on an old technique to give semantics to Algol-like languages. The basis has been extended to support modern concepts, such as object-oriented programming. Objects are represented by a box, or contour. Within the contour, the objects variables are shown, with name, type and value. The contour also uses arrows to point at other contours that are related, e.g. an other object representing the value of a variable, or an enumerator. Inheritance is shown by putting the contour of an object within the contour of the extended object. Object instances are kept separate from the contours representing inheritance, but will have relational links when necessary. Method calls are also represented in the diagram, in their own contours, linked to the calling object. JIVE offers to hide some of the information, such as inheritance or the composition of objects, in order to make the diagram smaller, and easier to read. Something that can be especially useful when working with larger programs, with many objects and relations. Visibility is aided by the use of colors to highlight specific elements of the diagram, such as variables bound to objects, and method-calls.
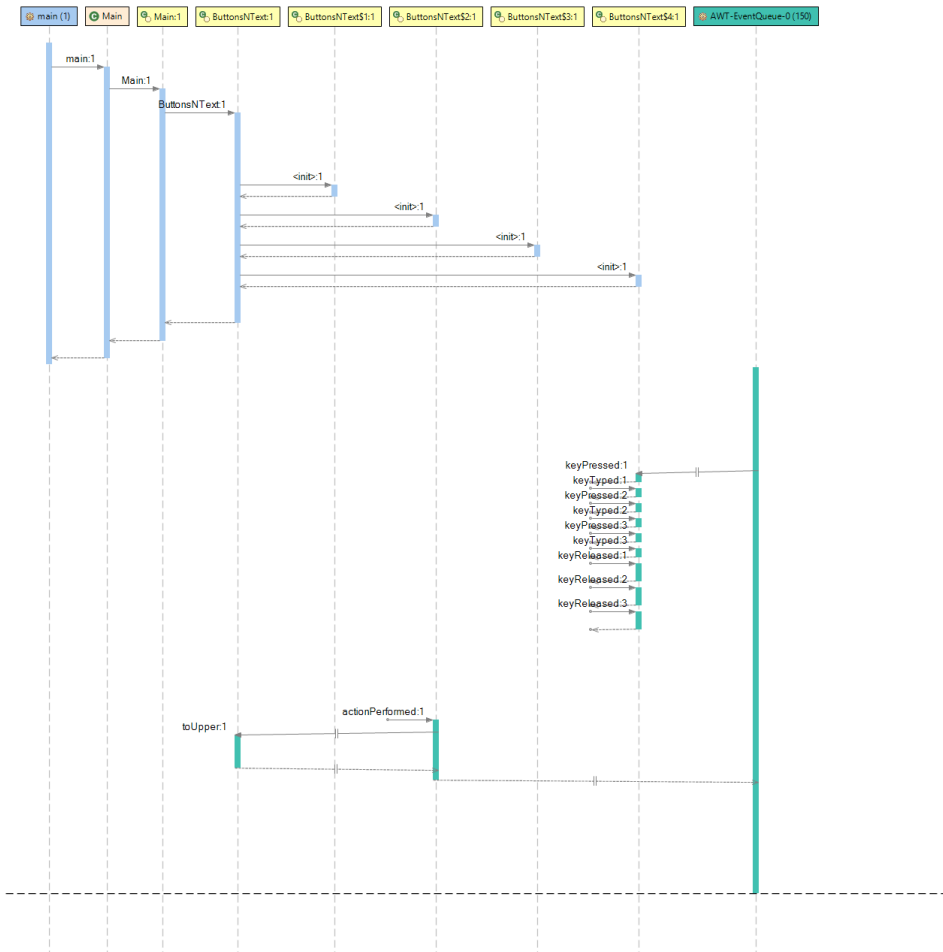
Figure 4.3: A sequence diagram generated by JIVE, while running an instance of MMI-Exercise 1

The sequence diagrams, shown in Figure 4.3, are fairly standard, with threads and objects represented by boxes on a row at the top, each with a vertical life-line. The actual sequence is shown with a thicker lifeline, with arrows representing method-calls. In order to differentiate the threads where the execution is happening, the sequences are colored with the same color as the thread- box the sequence originates from, regardless of witch objects and methods are involved. In Figure 4.3, one can clearly see the two colors representing the main- and the AWT-event-thread, and how elements in the diagram are colored by their parent thread.

Figure 4.4: The sequence diagram right click menu

Right-clicking on a lifeline offers the ability to collapse method-calls originating from that lifeline in order to hide unnecessary information, shown in Figure 4.4. The collapse-menu offers four options when used on the lifeline of an object: after, before, children and here. Before and after collapses lifelines that occurred before or after the selected event, at the same depth in the sequence-tree. Children collapses all events that are children of the selected event, while here collapses the selected event. Right clicking on an object instance at the top of the diagram also offers to collapse that objects entire lifeline, the result of which can be seen in Figure 4.5. While Figure 4.5 shows the result of collapsing the lifeline of an object, a similar result can be achieved by selecting 'collapse children' at the parent lifeline, or by selecting 'collapse here' on each of the three method-calls shown. The same options are of course available for expanding collapsed elements. Right-clicking also offers to set the execution-state via the jump-option, updating the contour-diagram and -model to show that state.

(a) Normal

(b) Collapsed

Figure 4.5: Normal and collapsed section of a sequence diagram

Both of the diagrams can be saved as a high-resolution image at any time, so that one can look at diagrams from earlier runs, instead of being forced to view them through JIVE. This also helps to visualize any changes made to a program,

and to see what the effects are on the program flow. The diagrams also share the ability to zoom in and out, further helping with the handling of larger diagrams.

Closely related to the diagrams, is the ability to quickly jump backwards and forwards between execution states.This is enabled by the trace-log, which contains an entry for every single event that occurs during the execution of a program, excluding those that are removed by the exclusion filter. Each event is assigned an identifying number, in ascending order, and information about thread, type, caller, target, and the location of the source-code is stored. The log is used as the basis for the models that make up the diagrams, and can be saved as both XML- and CSV-files for later use. As mentioned in the prestudy, logging every event has a significant impact on run-time performance, limiting the size and complexity of the programs that can be used with JIVE in a meaningful way. On the other hand, it allows the almost instant jumping between recorded states, as opposed to techniques that save a snapshot at predefined intervals, requiring the program to be run from the snapshot-state to the desired one, even if it is just a single step backwards.

In order to improve both performance and readability, the mentioned exclusion filter checks the origin of each event, and determines whether or not the event is to be logged. By default, the filter excludes the entire Java API, and by doing so, it is focusing on the classes that are made by the user.This filter can be adjusted to better suit the program being analyzed, by adding or removing entries in order to see more or less details of the execution.

The model-views each display an alternate view of their respective diagrams.They show the data-model representing the diagrams in a hierarchical structure, much like the organization of files and folders. Right-clicking on an event in the sequence model allows you to set the execution state to that event, and have the diagrams updated accordingly. Clicking on elements in the contour-model allows you to inspect the values of objects and their variables.

Finally, the trace-log enables the use of queries to search for specific events in the execution. The queries are presented through a new tab in the eclipse search-window as seen in Figure 4.6, easily accessible through the search-view, and comes with several pre-defined templates to simplify searching. For example, searching for when a certain variable gets a certain value, only requires the user to specify the variable-name, its parent, and the value.

Figure 4.6: The JIVE search panel

4.2 Shortcomings of JIVE

Even though JIVE offers several useful features, there is still room for improvements, both for general use, and more tailored towards development of graphical interfaces, as is the focus of the MMI-course.

The diagrams are possibly the most useful feature of JIVE, showing a lot of information in an understandable way, and offering some useful interaction. But there is still room for improvement: Inner types are displayed with their automatically generated, and fairly anonymous, name unless given a proper name in code. By default, most of the classes defined by the JRE itself are ignored, omitting potentially important information in GUI-applications. For the sequence-diagrams, it is naturally not important no show the internal behavior of every object, but they are also hidden in the contour-diagram. Related to this, is the lack of visual connections between standard-objects, and for example instances of listeners added to them. More related to the target group for this project, is a lack of differentiation between the different object-types in typical graphical architectures. Especially

MVC, whitch is a major focus of the MMI-course, has certain distinct types that are more important than others.

The sequence-diagrams can quickly become huge and hard to navigate, and the zooming function has very few levels outward, as well as leaving all text to small to read when zooming out. Additionally there does not seem to be any way to vertically compact the sequence-diagrams, making them unnecessarily long in cases where calls to ignored or hidden methods are being made. For instance when horizontally collapsing large parts of the diagram, leaving the parent lifeline at the same length as before, as can be seen in Figure 4.5b. Some of the papers on JIVE mention regex-folding,that could be used to substitute a series of events with a single new event, but it is still an experimental feature under development, and not available in the current release of JIVE.

Another potential for improvement is the process of stepping through recorded states, which currently requires manual interaction for each step, making complete replays of an execution infeasible due to the massive number of events recorded even for simple programs.

The search-view, while useful, is very strict in terms of what it looks for. For example, searching for the creation of an object requires its full name, including packages, so searching for creations of JButton-instances would require a search for javax.swing.JButton.

4.3   Suggestions for improvement

Based on the mentioned shortcomings, the following improvements are suggested.

The ability to detect an inner type with a generic name, and display its parent type in the diagrams instead of its own. This will make, among others, listener-heavy programs much more understandable as one can see what kind of listener each object is, instead of guessing, based on when it is invoked in the sequence-diagram. Further helping the same situation, is to visually link listeners to the objects they are added to, in the contour-diagram, making it clearly visible which object is being listened to by the different listeners. A crude visualization of these two changes, compared to the original, can be seen in Figure 4.7.

Finding ways to highlight the different parts of e.g. an MVC-architecture, making it clearly visible which objects make out the models, views and controllers, may further help the understanding of a program. But such highlighting must be balanced in order to not create a visual chaos of different colors. Adding symbols instead of colors might be a better solution, both maintaining the current color-scheme, and adding new information. The coloring, highlighting and naming of inner types should apply to both of the diagrams. The biggest challenge with

this kind of highlighting, is the process of detection. While listeners can easily be detected by their implementation of a listener-interface, detecting a controller or a model is not trivial. One can make assumptions based on class names, but they will be worthless if a program does not match those assumptions, and developers are not likely to change their programs to fit.

(a) Original

(b) Changed naming of inner types

(c) Listeners connected to the object listened to

Figure 4.7: Crude comparison of suggested changes to the contour diagram

Exploring changes to the default exclusion-filter, in order to provide more useful information out of the box, is also an option. This will provide a better experience for certain scenarios, but may be useless in other cases by providing too much unnecessary information. A way to easily switch filters by defining presets might be preferable. The filter might also be extended to support both exclusion and inclusion, allowing a more fine-grained selection of interesting classes. As an example, one might be interested in ignoring the entire javax-package, but still allowing javax.swing in order to see more GUI-related objects. It is also important to not allow too many packages through the filter, as the performance can suffer immensely from logging too many events.

Allowing more ways to interact with the diagrams, e.g. hiding elements or compressing sequences, should improve usability for larger programs and longer runs. The way the sequence diagrams currently work, object instances are added as they appear, resulting in later method calls to draw lines crossing large parts of the diagram, as shown in Figure 4.8.

Figure 4.8: A section of a larger sequence diagram, showing method calls crossing several unrelated lifelines

In the case of exploring events triggered by a listener, the ability to isolate the involved objects, and reorder the diagram could be very useful, providing a clean view of a smaller series of events.Using the same figure as an example, the rightmost, and longest, bar would be moved to the far left, the three second-longest would be in the middle, while the shortest ones would be to the left. In total, only the lifelines of those five objects would be visible, all the ones that are being crossed in Figure 4.8 are unnecessary, and would be hidden as shown in Figure 4.9. This function should be accessed by right clicking on the topmost event that is desired, and selecting a "view isolated" option. Whether the isolated view should appear in a new eclipse view, or simply modify the existing view of the sequence diagram, is something that should be determined by user-testing.
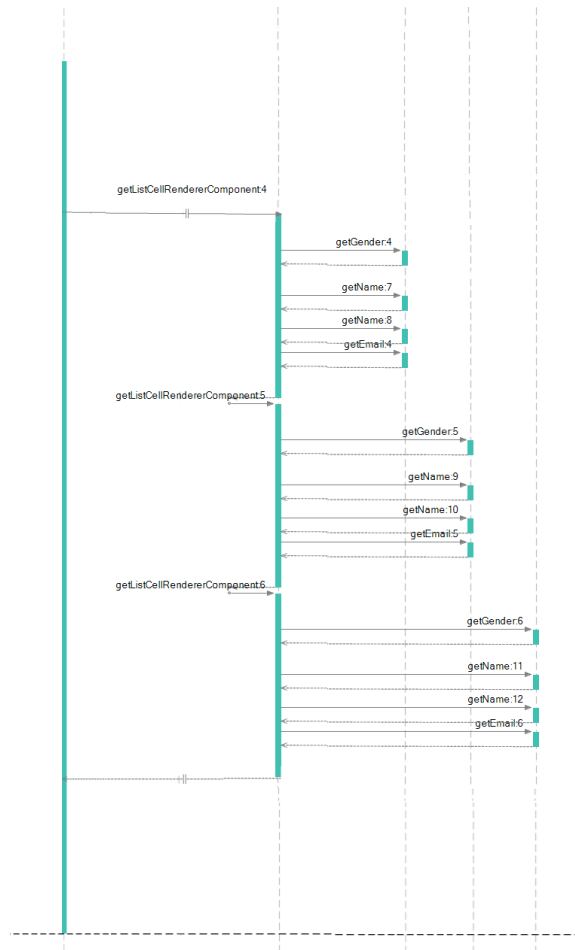
Figure 4.9: An example of how Figure 4.8 might look in the isolated view

Stepping through the recorded states is, as mentioned above, cumbersome. While there are quick and easy ways to jump straight to any interesting state, it may also be interesting to view a playback of a part of, or the entire execution. This would allow an easier way of observing changes happening in a program. The playback would automatically step through the recorded states at a pace that the user should be able to adjust, updating the diagrams for each step. Playback should be possible to initialize from any selected event that has been recorded, as starting from the beginning would often be a waste of time.

Searching can be improved by relaxing the requirements for for search-terms. The requirement of full class names in searches caused some confusion, and made

me wonder if the feature was working at all. Relaxing this requirement to allowing partial class names, or substrings in general, would be an improvement to the usability of the tool.

4.4   Implemented changes

The first change to be implemented was the identification and presentation of instantiated interfaces, which are now displayed in the diagrams with an appropriate icon, as well as being labeled with the interface it implements, instead of the generic class name that it is assigned by default. This function was also expanded to identify and label instances of abstract classes, and the lambda-expressions that were introduced with Java 8.In order to still allow the actual class name to be visible, the tool-tip shown when hovering the mouse pointer over an object remains unmodified, showing the actual class name and icon.

The filtering function was expanded with the ability to specify packages that are not to be excluded from the execution model, as suggested in section 4.3. Adding a package to be included is as simple as adding a '+' in front of the package-name when adding it to the filter. As an example, the default filter excludes the javax package, and any subpackages that are not specifically listed with a '+' in front. When adding the line '+javax.swing.*' in order to let swing components through the filter, one will let every subpackage of javax.swing though the filter as well, which is the intended design. Unfortunately, there are a lot of classes in both swing, and its subpackages that are not used directly when writing swing-programs, and that a user will have no interest in seeing in the diagrams. This can result in very poor performance, as unnecessary events are logged, and cluttered diagrams, but can be handled by adding the unwanted classes to the filter for exclusion. Depending on the package in question, the amount of extra items added to the filter can become quite large, and identifying all unwanted classes and subpackages may take hours at worst. One weakness in the filter, related to the problem of unwanted packages, is that due to its design, allowing the contents of a package, will also allow any classes that are contained directly within the parent package, as this also has to be removed from the internal list of exclusions that make up the filter. This is not a big problem in the case of 'javax.swing', as the 'javax' package does not contain any classes, but there are definitely other packages that will show this behavior.

With an appropriate filter, one will get to see the connection between listeners and the objects they listen to, via the existing object-containment-links. Such a filter would need to allow the object listened to, and the list it uses to organize its listeners, and is likely to become very large due to the extra exclusions one must add, as a consequence of letting something through. There is currently no easy way to quickly switch the entire filter, but by using multiple launch-configurations,

and modifying Eclipses launch file in a text editor, it is possible to distribute a pre-made filter, along with instructions, to at least avoid the repeated addition of single elements in the filter for each program being launched.

The isolated view was implemented as a separate view-tab, and does what was proposed in Figure 4.9: it displays the events caused by the selected event, and hides everything else. By right clicking on an event in the regular sequence diagram, and selecting the 'Isolated view' option, the isolated view is triggered.It is also possible to further focus on a part of the diagram from within the isolated view, by right clicking on the desired event, and selecting the 'Isolated view' option again. All of the functionality from the regular sequence diagram has been retained in the isolated view, so that the only difference is what parts of the execution are visible.

The searching functionality was relaxed by allowing partial matches, and disabling case sensitivity. While relaxing the search may cause false positives, a user should be able to identify such cases quickly, or at least identify the results they were looking for. The fact that partial matches and case insensitivity is the standard behavior in most search-engines, sets an expectation for the behavior of the search within JIVE. Unfortunately, due to the existing implementation consisting of separate searching- and matching-methods for each search type, not all searches were updated to this relaxed state.

4.5   Evaluating usefulness

In order to evaluate both the usefulness of JIVE, and the changes that were implemented, testing was performed with a group of students. During this test, the students received a demonstration of what JIVE is, before getting to use it on an implementation of the first, and the last exercise of the MMI-course. During the test, the students were asked to evaluate the perceived performance, and whether it would affect the desire to use JIVE. The actual performance was measured separately from the user testing, the test was merely to see if this was an acceptable trade-off. They were also tasked with evaluating to what degree they found JIVE to be helpful in understanding the example programs, and how understandable the diagrams were. Lastly, they were asked to evaluate the usefulness of the modifications detailed in the previous section, and whether or not they did what was intended.

Regarding performance, it has already been stated that the tracing behavior of JIVE must affect the analyzed program. The specific penalty varies with the complexity of the program, and the level of detail that is being logged. With the default filters, the programs used in the test had an average startup time of INSERT TIME, and responded to user input with a delay of no more than one second. With the modified filter mechanism, and filters adjusted to include information on relevant classes, the startup time increased to INSERT TIME2, while the user input delay ended up at RESPONSETIME HERE.

Chapter 5

Conclusion

This section will summarize the previous sections in light of the research questions proposed in the introduction.

5.1 Research Questions

Research Question 1 *What is the current state of the various visualization tools that are available?*

There are currently several different tools that can aid a developer in his understanding of programming, ranging from the regular debuggers found in most IDEs, to code analysis and execution tracers. Most of these focus on one or two features, instead of attempting to combine them all. Regarding NTNUs chosen teaching environment, the amount of available tools is naturally reduced, but there is still a large portion of tools available, thanks to the plug-in support of Eclipse. Among these, JIVE was found to be of most interest, due to its combination of features, including tracing, diagram generation, and state-jumping.

Research Question 2 *Could any of these be integrated into the current teaching environment at NTNU, consisting of Java and Eclipse?*

Both JIVE and other tools are available as Eclipse-plugins, and as such, they can easily be integrated into the courses that use Eclipse. Unless a very platform-specific mandatory framework is used, students are of course free to choose a different IDE while following a course, and by doing so, they will also opt out of the use of JIVE or any other tools. But these students are not likely to be the ones that would benefit the most from such tools in the first place, as a decision to not use recommended tools usually indicates that they know what they are doing. As mentioned in [Gestwicki and Jayaraman, 2005, p. 99], JIVE has already been successfully used to explain design patterns to students in a graduate-level seminar, indicating that it is suitable as an educational aid.

Research Question 3 *Is there room for improvement in how these tools are used, and the ease of using them?*

As detailed in chapter 4, there are definitely a potential for improvements in how JIVE works, and is used. While the existing features are useful as they are, they were found to be lacking in usability. The diagrams were found to lacking in their description of certain classes, making identification difficult. Larger programs were also found to quickly grow the sequence diagram to large sizes, making it harder to get an overview.

Some of these points were the focus of the improvements that were implemented in section 4.4, and their results are summarized when answering the last research question.

Research Question 4 *Would the use of such tools and any improvements actually be useful for the students?*

5.2   Future work

keywords:
If gui is a desired use case: Focus on performance, minimize unnecessary work/avoid work as soon as possible.
JFX is essentially not usable with JIVE in current state: too many things going on behind the scene

in general:
further improvements on usability.
compression of diagram height
looser search terms
improving filter, tree-view of classes/packages found in workspace with checkboxes for inclusion/exclusion?

Glossary

Breakpoint A source code marker telling the debugger to halt program execution at a certain point. 5

Code Canvas A visualization-tool for Microsoft visual studio, showing code, diagrams and documents on a large layered canvas.
http://research.microsoft.com/en-us/projects/codecanvas/ 6

Command Line Interface A text-based interface for interacting with programs via a terminal. 6, 26

Contour diagram An enhanced object diagram, showing objects, their variables and their relations to other objects.
[Jayaraman and Baltus, 1996, Streib and Soma, 2010] 7

Debug Visualization Plugin An Eclipse plugin that provides a graphical view of variables during debugging.
https://code.google.com/a/eclipselabs.org/p/debugvisualisation/ 6, 7

Execution trace A log of all changes to the state of a program throughout its execution. 5

GNU debugger A multiplatform, multilanguage CLI-debugger with tracing.
http://www.sourceware.org/gdb/ 6, 26

Integrated Development Environment A software application that provides facilities for software development such as source code editor, compiler etc. 2, 26

JAVAVIS Tool that generates UML diagrams from running java applications.
[Oechsle and Schmitt, 2002] 7

Jinsight An advanced debugger made by IBM, supports visualization, and powerful analysis.
[Pauw and Vlissides, 1998] 7

JIVE An advanced debugging tool supporting visualisation, backward stepping, and querying.
http://www.cse.buffalo.edu/jive/ 7

Model View Controller A design pattern defining an architecture where data is kept separate from the component that displays it to a user. 2

Trace Viewer Plugin Tool to visualize and analyze communication of parallel message passing programs.
[Kranzlmüller and Klausecker, 2009] 6, 7

Trace-Oriented Debugger Trace-Oriented Debugger. A debugging tool that executes queries on program traces.
[Pothier et al., 2007] 6, 7

Whyline A query-based debugger that provides an easy way to find out why things are as they are.
[Ko and Myers, 2009] 6

Acronyms

CLI Command Line Interface 6

GDB GNU debugger 6

IDE Integrated Development Environment 2, 4, 6

MVC Model View Controller 2, 3

Bibliography

[Gestwicki and Jayaraman, 2005] Gestwicki, P. and Jayaraman, B. (2005). Methodology and architecture of JIVE. *Proceedings of the 2005 ACM symposium on ...*, 1(212):95–104.

[Jayaraman and Baltus, 1996] Jayaraman, B. and Baltus, C. (1996). Visualizing program execution. *Proceedings 1996 IEEE Symposium on Visual Languages*, pages 30–37.

[Ko and Myers, 2006] Ko, A. and Myers, B. (2006). An exploratory study of how developers seek, relate, and collect relevant information during software maintenance tasks. *Software Engineering, ...*, 32(12):971–987.

[Ko and Myers, 2009] Ko, A. J. and Myers, B. a. (2009). Finding causes of program output with the Java Whyline. *Proceedings of the 27th international conference on Human factors in computing systems - CHI 09*, page 1569.

[Kranzlmüller and Klausecker, 2009] Kranzlmüller, D. and Klausecker, C. (2009). Scalable Parallel Debugging with.

[Oechsle and Schmitt, 2002] Oechsle, R. and Schmitt, T. (2002). Javavis: Automatic program visualization with object and sequence diagrams using the java debug interface (jdi). *Software Visualization*, pages 176–190.

[Pauw and Vlissides, 1998] Pauw, W. D. and Vlissides, J. (1998). Visualizing object-oriented programs with jinsight. *Object-Oriented Technology: ECOOP'98 ...*, pages 541–542.

[Pothier et al., 2007] Pothier, G., Tanter, E., and Piquer, J. (2007). Scalable omniscient debugging. *ACM SIGPLAN Notices*, 42(10):535.

[Streib and Soma, 2010] Streib, J. and Soma, T. (2010). Using contour diagrams and jive to illustrate object-oriented semantics in the java programming language. *Proceedings of the 41st ACM technical symposium ...*, pages 510–514.